

Software and Hardware Techniques for SEU Detection in IP Processors

C. Bolchini · A. Miele · M. Rebaudengo · F. Salice ·
D. Sciuto · L. Sterpone · M. Violante

Received: 27 November 2006 / Accepted: 11 July 2007 / Published online: 29 January 2008
© Springer Science + Business Media, LLC 2007

Abstract In the recent years both software and hardware techniques have been adopted to carry out reliable designs, aimed at autonomously detecting the occurrence of faults, to allow discarding erroneous data and possibly performing the recovery of the system. The aim of this paper is the introduction of a combined use of software and hardware approaches to achieve a complete fault coverage in generic IP processors, with respect to SEU faults. Software techniques are preferably adopted to reduce the necessity and costs of modifying the processor architecture; since a complete fault coverage cannot be achieved, partial hardware redundancy techniques are then introduced to deal with the remaining, not covered, faults. The paper presents the methodological approach adopted to achieve the complete

fault coverage, the proposed resulting architecture, and the experimental results gathered from the analysis of the fault injection campaigns.

Keywords Reliability · Hardware/software techniques · Single-event upset faults · Fault injection

1 Introduction

In the most recent years, given the relevant computational power of IP cores, the design of reliable systems by means of software techniques received a lot of attention, due to the interesting possibility to use a commercial processor core, without requiring any customization. Indeed, such an advantage is very important when considering time to market, the flexibility of the software, as well as the independence of the system from a specific IP core vendor. Nevertheless, the drawback of a pure software approach is the impossibility of achieving a complete fault coverage [4], due to the presence of a limited, but not empty, set of faults that may cause the processor not to behave correctly, without detecting the error. As a consequence, in safety-critical systems, where it is mandatory to be able to guarantee and assess a complete fault coverage, a combined software-hardware approach needs to be adopted. Our approach aims at introducing a hardware fault detection mechanism for those errors that may not be detected by means of software techniques; costs and benefits will be evaluated, achieving the desired complete fault coverage with respect to the considered the Single Event Upset (SEU) fault model.

It is worth noting that this approach is not only the result of the systematic analysis of software techniques for hardware fault detection and tolerance, but it is also a mixed hardware/software co-design solution for the design

Responsible Editor: N. A. Toubia

C. Bolchini (✉) · A. Miele · F. Salice · D. Sciuto
Dip. di Elettronica e Informazione, Politecnico di Milano,
Milan, Italy
e-mail: bolchini@elet.polimi.it

A. Miele
e-mail: miele@elet.polimi.it

F. Salice
e-mail: salice@elet.polimi.it

D. Sciuto
e-mail: sciuto@elet.polimi.it

M. Rebaudengo · L. Sterpone · M. Violante
Dip. di Automatica e Informatica, Politecnico di Torino,
Turin, Italy

M. Rebaudengo
e-mail: reba@polito.it

L. Sterpone
e-mail: sterpone@polito.it

M. Violante
e-mail: violante@polito.it

of reliable systems. In fact, the information on fault coverage derived from the qualitative, methodological approach allows an estimation of costs and benefits of the pure software techniques, whereas overheads and fault coverage for classical hardware techniques are well-known. Therefore, in a reliability-aware hardware/software co-design flow [3] where a complete fault coverage is required and flexibility and area constraints drive the solution space exploration, the achieved architectural result is a valid solution. We can consider the 100% hardware redundancy technique as the solution at one end of the solution space providing a complete fault coverage. On the other end of the solution space, with a complete fault coverage, there is the application of all available software techniques, supported also by a partial hardware redundancy to cover all faults, identified as *all-software and partial hardware* approach. Other solutions lie in between, with a different mix of software and hardware techniques and still providing a complete fault coverage, characterized by various overheads in terms of area, code size, performance degradation, etc.; in this paper we address the results for the boundary solutions preliminarily presented in [6] and we explore the solution space by identifying the most promising nonredundant software/hardware mix of techniques.

The main contribution of this paper is a revised version of the preliminary qualitative analysis of soft errors presented in [4], and a modified application of redundancy techniques with respect to those presented in [2], driven from the results of the presented analysis. As experimental results will show, the methodological approach enabled a more specific hardware customization, overcoming the limitation of the previous work. An architecture-independent methodology is presented, whereas the experimental phase refers to the Leon2 processor architecture [9].

The rest of the paper is organized as follows. Section 2 briefly introduces the adopted fault model and presents the proposed combined approach designed to cover all possible errors caused by SEUs, discussing the software and the hardware techniques used to pursue the desired goal. The resulting solution has been validated from a theoretical point of view, nevertheless a fault injection campaign has been also carried out, to further verify that all faults are covered and new ones are not introduced; the results are reported in Section 3. Final remarks and future work close the presentation.

2 The Proposed Approach

The approach here presented, starts from the conclusions drawn in [2] and proposes an enhancement based on a more systematic analysis of the fault/error relationship. The methodology combines software and hardware techniques

for fault detection to introduce reliability properties into software modules of critical embedded systems. The methodology is mainly based on the software approach due to its benefits—acceptable error coverage and reduced design costs and overhead, coupled with a performance degradation and code size growth—and exploits a partial hardware replication to provide the detection of faults that are not covered by the software techniques. The adopted fault model is introduced in the following, then an evaluation of the software techniques will be presented. The goal is the identification of the areas of the processor whose functionalities are not completely covered, leading to the introduction of partial hardware replication.

2.1 Fault Model

Single programmable logic devices, such as System-on-Chip (SoC) and Field Programmable Gate Arrays (FPGAs), are today commonly used to implement embedded systems. Such devices are susceptible to errors resulting from the perturbation of memory cells, provoked by ionizing radiation ([8]). An example of this effect is the Single Event Upset (SEU), i.e., an undesired change of state in a bistable or storage element, also referred to as a soft error. SEUs are a major cause of concern in a space environment, but they have also been observed at ground level [13]. In a combinational circuit, e.g., an arithmetic logic unit (ALU), ionizing radiation may induce a transient voltage pulse that can lead to incorrect output. The modification of the memory cell is not permanent, since a successive write operation can override the erroneous content and replace it with a new one. The proposed approach aims at providing a complete fault coverage for any single SEU in the storage elements of a processor, in a SoC or in a FPGA, occurring within the execution of an instruction and thus having the unique effect of causing a bit-flip in one of the storage elements. For the analysis of the software techniques, the considered fault is modeled at an higher abstraction level, as further detailed in the next subsection.

2.2 Software Techniques

Software techniques for fault detection are a well explored issue in literature ([1, 5, 10, 11]). The basic idea is to apply temporal and information redundancy strategies on the program code, modified by introducing additional instructions and variables to perform equivalent computations more than once and to compare results. Three main techniques are considered:

- Data processing instruction replication ([5, 11]). Replication of all the variables of the program and all the data processing instructions. This replication aims at

performing redundant data elaborations to check redundant results.

- Self checking block signature ([1, 10]). Static control of the execution flow; the technique assigns to each basic block of the program (the maximum branch-free block of instructions) a static signature and introduces, at the beginning and at the end of each basic block, two instructions for checking and updating the runtime signature stored in a global variable against the static signatures.
- Conditional branch instruction replication ([11]). Replication of the branch instruction in each branch direction, aiming at performing a run-time check of the direction taken during the execution of a branch instruction.

This set of techniques is, though, not sufficient to cover all possible errors caused by SEU faults, as it will be shown in the following. The first attempt to classify such open issues has been adopted as the starting point ([4]), to carry out an architecture-independent analysis, able to fit different processors. For this purpose, the adopted architecture is described in a behavioral fashion, to maintain a high level of abstraction and to describe the essential aspects of any processor without considering structural aspects that are peculiar to the implementation of one processor. The model is composed of a *program executor* and two memories (a *data memory* and an *instruction memory*). The program executor is a high level description of a single processor architecture, modeled only with its workflow, shown in Fig. 1, composed of several basic operations that are executed sequentially step-by-step, for each instruction of the program code. While considering the imperative programming paradigm, the instructions that the model executes can be classified into two classes: *data processing instructions* and *control instructions*. The former are executed sequentially starting from the entry point of the program, while the latter, only, can modify the instruction flow by performing a jump to any other instruction.

When considering the adopted architecture, the fault may be modeled as a corruption of the execution of a single step of the workflow or as a corruption of the content of one of

the two memories. When considering the *program executor*, the result of the physical fault at the behavioral abstraction level is represented as a corruption of a value, and the effect is an erroneous execution of a step of the processor workflow. If the fault corrupts the memories, an error is generated and is detectable when the corrupted value is used for a computation. When an error affects the program execution, the data or the control flow of the program is corrupted. When a value used for the data elaborations is corrupted, an error is generated in the data flow. When an error affects the control flow, it causes an erroneous evolution of the execution flow; in particular, the effect is a wrong jump performed within the instruction code and, thus, the execution continues with a wrong instruction. More in detail, errors affecting the control flow can be classified as follows:

- during the execution of an instruction inside a basic block, instead of the current instruction, a jump may be performed to another instruction either (a) in the same basic block, or (b) in a different basic block.
- during the execution of the control instruction at the end of a basic block, a jump may be performed either (a) to the beginning of a wrong basic block, or (b) to the middle of a basic block.

Given the results of the analysis of the possible faults in the architecture and of the corresponding errors affecting the program execution, it is possible to study the cause/effect relationship between faults in the architecture and errors in the program execution. When considering the adopted program executor, it is possible to state that both the functionality corrupted by the fault and the type of instruction that the processor is executing are responsible for the type of generated error and its effects on the program execution. The fault/error relationship is shown in Tables 1 and 2, where a specialization of the above listed errors is used: when executing a data processing instruction, a fault generally causes an error in the data flow; in case the computation of the next instruction is corrupted or the current instruction is transformed into a control instruction, an error may be generated also in the control flow. On the other hand, when executing a control instruction, the effect of the fault is generally an erroneous evolution of the control flow.

When the fault affects one of the two memory elements the error will manifest its effects when the corrupted value is used by the executor. If the data memory is affected, the error will corrupt the data flow if the value is used during the execution of a data processing instruction; otherwise the control flow is affected. When the instruction memory is corrupted, an instruction is transformed into a new one: therefore, the error will affect the data flow if a data processing instruction is transformed into another data processing instruction; otherwise also the control flow is affected.

Fig. 1 Program executor workflow

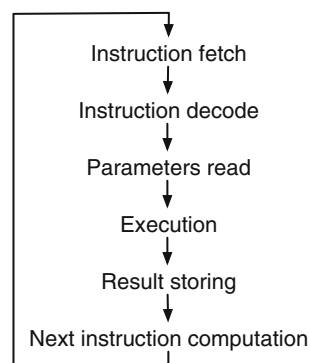


Table 1 Errors possibly caused by faults affecting a data processing instruction

Corrupted phase	Generated errors
<ul style="list-style-type: none"> ○ Instruction fetch ○ Instruction decode 	<ul style="list-style-type: none"> ○ A data processing instruction is transformed into: <ul style="list-style-type: none"> □ another data processing instruction a wrong result is produced □ a control instruction: an erroneous jump to a random target is performed <ul style="list-style-type: none"> ▷ in the same basic block, or ▷ in a different basic block
<ul style="list-style-type: none"> ○ Parameters read ○ Execution ○ Result storing 	<ul style="list-style-type: none"> ○ Instruction execution leads to a wrong result that is stored in memory
<ul style="list-style-type: none"> ○ Next instruction computation 	<ul style="list-style-type: none"> ○ An erroneous jump to a random instruction occurs: <ul style="list-style-type: none"> □ in the same basic block, or □ in a different basic block

The presented fault/error relationship provides the starting point to perform a critical analysis of the most interesting software techniques for fault detection. The purpose is to identify the errors each technique covers and the errors not covered by any technique, to determine the functionalities of the processor not completely dealt with by the considered software techniques.

The qualitative analysis shows that, by means of the adopted techniques, a good capability of error detection is achieved for the data flow of the program. In fact, due to the replication of all the data processing instructions and of all the variables, the fault affecting the processor will corrupt a value and the checking instruction will detect an incongruence between two redundant values. On the other hand, the study shows that the techniques present some limitations in the detection of errors that affect the control flow. Even if the techniques of block signature and conditional branch instruction replication make it possible to identify erroneous jumps among different basic blocks and wrong evaluations of branch directions, a partial error detection capability is obtained in case of erroneous jumps to instructions of the

same basic block. In fact, the block signature technique is not able to detect this kind of error; moreover, the replication of the data processing instructions makes it possible to detect this kind of error only if the erroneous jump generates an incongruence in the data flow, i.e., if only one of the two replicas of an instruction is executed.

Another situation that cannot be covered by software techniques is related to the control transfer performed when an exception is thrown due to a trap or an interrupt. In this scenario the block signature technique cannot be applied because the jump is not performed by an explicit instruction but is requested by an interrupt or a trap. It is worth noting that this scenario is not taken into account for this study because no specific operating system is considered, and the management of these issues strictly depends on the adopted operating system.

The results of the analysis show that the functionalities not covered by the software techniques are those that cause an erroneous jump to an instruction of the same basic block. The fault that may originate this error is in the next instruction address computation. Moreover, the decoding phase may be involved in scenarios that lead to the same

Table 2 Errors possibly caused by faults affecting a control instruction

Corrupted phase	Generated errors
<ul style="list-style-type: none"> ○ Instruction fetch ○ Instruction decode 	<ul style="list-style-type: none"> ○ The control instruction is transformed into: <ul style="list-style-type: none"> □ a data processing instruction: a wrong result is produced and no jump is performed □ another control instruction: an erroneous jump to a random target is performed <ul style="list-style-type: none"> ▷ in the same basic block, or ▷ in a different basic block
<ul style="list-style-type: none"> ○ Parameters read ○ Execution 	<ul style="list-style-type: none"> ○ Jump target computation error: an erroneous jump to a random instruction occurs: <ul style="list-style-type: none"> □ in the same basic block, or □ in a different basic block
<ul style="list-style-type: none"> ○ Next instruction computation 	<ul style="list-style-type: none"> ○ Branch direction evaluation error: wrong branch direction is taken ○ The execution of a jump instruction is affected: no jump is performed
<ul style="list-style-type: none"> ○ Result storing 	<ul style="list-style-type: none"> ○ No instruction is performed: the fault causes no error because it is not activated

kind of error: in case an instruction is transformed by a fault into a control instruction, a random jump may be performed to another instruction of the same basic block. Finally, this kind of error may be generated by faults affecting the checking instruction introduced by the replication of the data processing instructions; in fact, when this technique is applied, a checking instruction is introduced into the code and it is not possible to control its execution by means of the block signature. Thus, it is possible to state that a part of the processor functionalities is not completely covered by means of the software approach. In the next subsection it will be shown how to overcome the highlighted limitations.

2.3 Hardware Techniques

By referring to an architecture model, the adopted program executor may consist of several modules that implement the described basic functionalities. Therefore, it consists of a control module composed by units such as an instruction fetch unit or an instruction decode unit, and a data-path composed by units such as an ALU or memory access interfaces. By considering the architecture model, it is possible to identify the units whose functionalities are not completely covered by the software techniques; such units are the program counter management unit, the decoding unit and the branch management unit.

When considering the features of the adopted fault model—the soft error usually affects storage elements—it is necessary to act on the registers of the uncovered units by means of hardware replication and comparison. Therefore, the subset of registers to be replicated depends on the structure of the units that implement the selected functionalities. Such set of elements includes the program counter and the instruction register; according to the processor structure, other control registers should be duplicated, for instance the registers used for transmitting the jump address between the several units. Moreover, when considering a pipelined processor, several pipeline stage registers containing information related to the instruction word, the program counter and the jump address should be replicated.

The partial hardware replication here applied consists in duplicating the identified subset of registers; for each one of them, a redundant register is added. The basic register is used for the nominal computation and the redundant one is used for storing a copy of the value. A checker is used to compare the contents of the two replicas: when inconsistencies are observed, the system can be held and reset to prevent wrong values from being propagated. It is worth noting that the adoption of this strategy of hardware replication does not affect processor performance and guarantees concurrent error detection without additional computational overhead, although it implies an overhead in terms of area.

2.4 Software Tuning after Hardware Redundancy Introduction

As discussed above, the impact of each single technique in terms of the provided fault coverage is not completely disjointed, especially when software methods are related to hardware techniques. For example, the Self-Checking Block Signature method (software bound) allows the identification of a subset of SEUs modifying the program counter, whereas the hardware duplication of the program counter (including the registers which span through the pipeline) allows the identification of any SEU affecting the program counter.

As a consequence, the necessary application of hardware techniques to achieve the complete fault coverage may turn one or more of the software techniques redundant; in fact, if the duplication of a register was introduced because software redundancy could cover only a subset of its faults, then keeping the software technique provides no benefit at all. Therefore, in the design methodology, the application of the hardware redundancy techniques must be followed by an additional step to reconsider the benefits provided by each software technique, in order to discard those which have become redundant from the fault coverage point of view, that nevertheless introduce performance degradation and memory costs.

It is important to remember that the initial goal is to work with software techniques, to benefit from the flexibility it offers, in allowing the designer to possibly decide the level of fault coverage to be pursued, and the performance degradation that is considered acceptable. Indeed, if a complete fault coverage is desired, hardware modifications are necessary and thus, to limit the introduced overhead (code size, performance degradation, area overhead, etc.), software techniques are then discarded if the faults they cover are already detected by the hardware techniques.

The resulting design methodology consists of four steps, to be performed in sequence in order to achieve the final implementation with the desired fault coverage and with the minimal set of necessary and sufficient software and hardware techniques (Fig. 2).

In the first step, a set of software techniques is introduced by selecting, among the available techniques, either the most promising ones or all the available techniques. On the initial hardware architecture where the obtained enhanced software application runs, a fault coverage analysis is performed in order to identify both the uncovered faults and the involved hardware parts (step 2). In step 3, a set of necessary and sufficient hardware techniques is implemented; in particular, a set of modifications on the involved architectural parts (e.g., the CPU) is applied in order to achieve a 100% fault-coverage of the hardware/software architecture. Finally, in step 4, the irredundant final solution is identified by eliminating the redundant subset among the applied software techniques (e.g., software techniques covered by the intro-

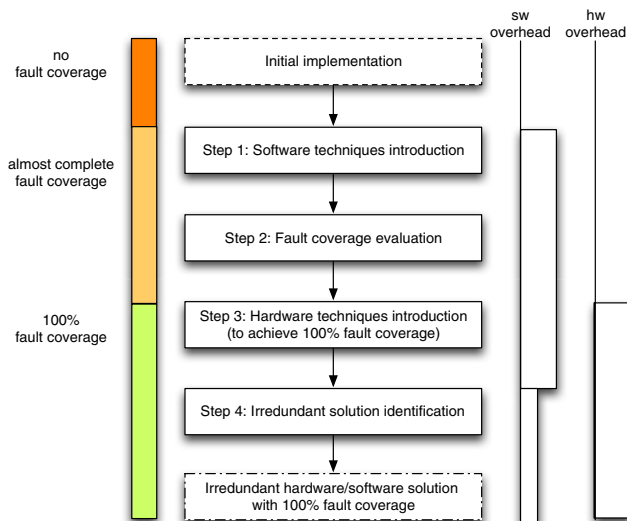


Fig. 2 100% fault coverage design methodology

duced hardware ones); the final result is an optimal software/hardware solution with 100% fault coverage.

3 Methodology Evaluation

In order to evaluate the proposed hardening methodology we analyzed a case study based on a microprocessor system

considering a fault injection system based on an emulation environment implemented on a SRAM-based FPGA board [7]. This environment is able to inject SEUs within all the memory elements of a microprocessor core. The evaluation flow we adopted consists first in hardening the microprocessor system considering the proposed Program Executor workflow model. Then, several fault injection campaigns aimed at demonstrating the effectiveness of the proposed hardening methodology are performed.

3.1 Case Study

The case study we adopted is the LEON-II microprocessor core [9], a processor based on an integer unit (IU) that implements SPARC integer instructions. It has several features such as a five-stage instruction pipeline, separate instruction and data cache interface and an architectural support for 2 to 32 register windows. A scheme of the LEON-2 integer unit is shown in Fig. 3, where the modules of the integer unit are organized in a five-stage pipeline. The proposed program executor workflow is adaptable to the considered architecture, where the program executor's phases model the pipeline stages and the program executor's modules implementing the basic functionalities model the registers and the functionality units. In details, the instruction

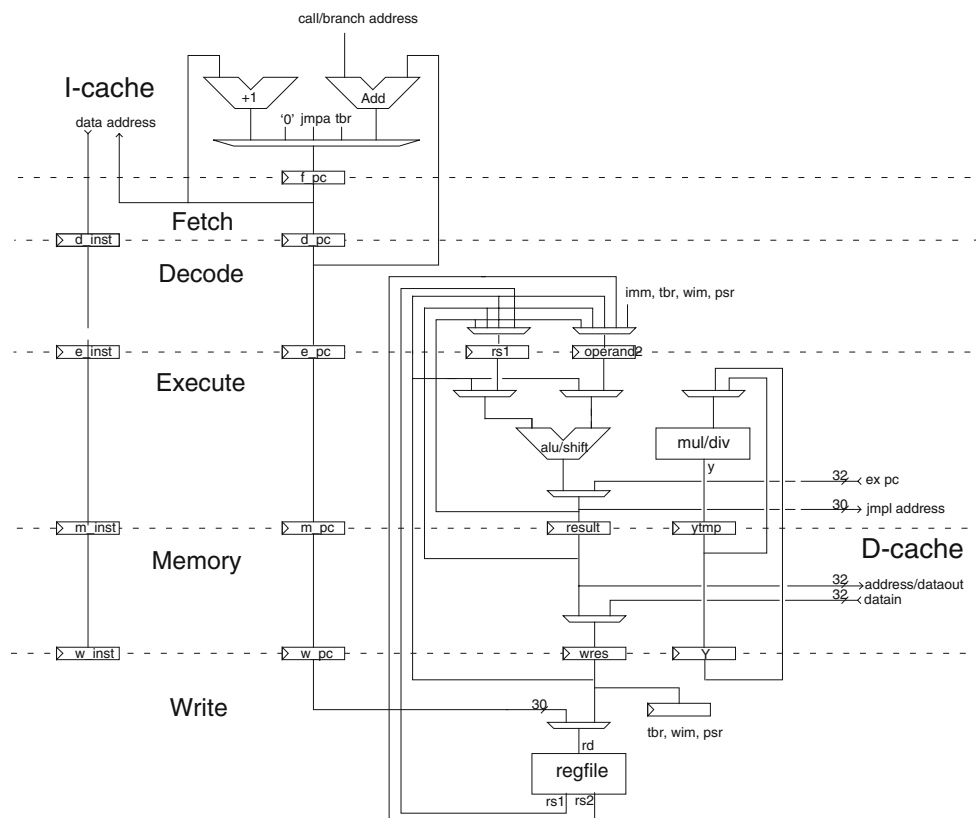


Fig. 3 LEON-2 integer unit block diagram

Table 3 Experimental results: for each benchmark and implementation strategy we report the number of undetected faults per pipeline stage, and the implementation costs in terms of time and memory (times are in clock cycles and code size expressed in bytes)

Program	Version	Pipeline stage					Total	Exec time	Data size	Code size
		FE	DE	EX	ME	WR				
FIR	plain	190	142	67	39	11	449	6235	200	860
	sw(+hw)	3	0	0	0	0	3	29293	400	2864
EWF	plain	210	110	69	36	41	466	9854	164	724
	sw(+hw)	3	0	0	0	0	3	110955	328	4408
KLMN	plain	220	157	99	64	16	556	32811	6472	1216
	sw(+hw)	2	0	0	0	0	2	258445	12944	6184

fetch stage corresponds to the pipeline fetch, the instruction decode and the parameters read model the decode stage, the execution stage models the pipeline execute, finally the result storing and the next instruction computation stages model the memory and the write states, respectively.

We applied the analysis methodology described in Section 2.2 to the LEON-2 processor, and we identified those components that, in the considered processor, implement the program counter management unit, the decoding unit and the branch management unit. In order to detect those occurrences of SEUs in these units not covered by software rules, the following registers must be hardened: the branch address register, the jump address register, the program counter, the instruction registers, the trap control register, the result register, the annul register (used to flush the pipeline in case of hazard), and the register controlling the execution of instructions that span over more than one clock cycle. For the sake of this work, we adopted a simple duplicate-and-compare approach: we duplicated the registers, and we inserted a checker that signals the occurrence of any SEUs affecting them. The adopted solution does not affect the performance of the processor, while it introduces a limited area overhead, which can be estimated to be about 2%.

3.2 Experimental Results

To assess the effectiveness of the proposed approach we exploited the environment developed in [7] that provides the required accuracy since it allows accessing all the memory elements the processor embeds with a suitable time resolution. The environment consists of a FPGA board used to emulate an instrumented model of the considered processor, and of a Personal Computer that hosts the FPGA board and runs the software managing the fault injection experiments. The model of the processor is enhanced (i.e., instrumented) with suitable hardware components allowing observing and controlling each memory element, through which SEUs are injected, and SEU effects are observed. The fault injection experiments have been performed considering 3 benchmark programs:

1. Finite Impulse Response filter (FIR): A filtering of 16 fixed-point samples.

2. Fifth Order Elliptical Wave Filter (EWF): An elliptic filter over a set of 32 fixed-point samples.
3. Kalman filter (KLMN): A Kalman filter on a set of 16 fixed-point samples.

For the considered benchmark applications, we compared 3 implementations strategies:

- a. Plain: It is the plain version of the considered application. Hardware and software hardening techniques are not applied.
- b. All-software: It is the software hardened version of the considered application. This is obtained using only software techniques as it is explained in Section 2.2.
- c. All-software and partial hardware: It is the hardened version of the application, obtained using the combination of software and hardware techniques previously described.

We analyzed the dependability of the three implementations by performing the injection of 1,000 randomly selected faults in each of the five pipeline stages that compose the considered processor: the fetch stage (FE), the decode stage (DE), the execution stage (EX), the memory-access stage (ME), and the write-back stage (WE). The attained results are reported in Tables 3 and 4,

Table 4 Experimental results: for each benchmark and implementation strategy we report the overheads

Program	Version	Exec time ratio	Data size ratio	Code size ratio	Memory ratio
FIR	plain	1.00	1.00	1.00	1.00
	sw	4.70	2.00	3.33	3.08
	sw+hw	4.70	2.00	3.33	3.08
EWF	plain	1.00	1.00	1.00	1.00
	sw	11.26	2.00	6.09	5.33
	sw+hw	11.26	2.00	6.09	5.33
KLMN	plain	1.00	1.00	1.00	1.00
	sw	7.88	2.00	5.09	2.49
	sw+hw	7.88	2.00	5.09	2.49

Table 5 Experimental results for the nonredundant sw/hw implementation with a complete fault coverage

Program	Version	Pipeline stage					Total	Exec time	Data size	Code size
		FE	DE	EX	ME	WR				
FIR	data+br+hw	0	0	0	0	0	0	24589	400	2456
	data+hw	0	1	2	2	0	5	21418	400	2212
EWF	data+br+hw	0	0	0	0	0	0	107032	328	4264
	data+hw	0	1	2	1	0	4	104917	328	4088
KLMN	data+br+hw	0	0	0	0	0	0	207240	12944	4872
	data+hw	0	1	1	3	0	5	150572	12944	3816

where the number of SEUs provoking the program to compute wrong results are listed, along with the execution time, the data size, and the code size (these values are identical in both versions, software techniques only, and software and hardware techniques, thus we only reported them once). It can be observed that about 10% of the injected faults lead the plain version of the programs to produce wrong results. The considered faults are thus relevant, and effective countermeasures are needed. The programs hardened according to the all-software approach show promising results in terms of fault detection, we indeed observed a reduction of two orders of magnitude in the number of wrong results. However, few faults still escape the software detection mechanisms, confirming the necessity to introduce a further hardening mechanism. The results gathered when considering the combined all-software and partial hardware approach show that the number of faults causing wrong results is reduced to zero, confirming the analysis described in Section 2. These experiments outline the effectiveness of integrating techniques based on software modifications, with techniques that require limited hardware modifications. However, the achieved improvement of system's dependability comes at a cost of increased execution time (due to the time redundancy and the additional instruction for checking purposes we inserted in the software), memory overhead (due to the replicated data and instructions), and area overhead (due to the replicated registers). For the considered benchmarks, the average increase of the execution time is about six times, the memory overhead is about four times, while the area overhead is about 2%.

In order to find the best mix of software techniques to be adopted in hardening the source code of the programs after the introduction of hardware techniques, a second set of experiments has been performed (according to the Step 4 reported in the design flow depicted in Fig. 2). We analyzed the benchmark implementation where the all-software and partial hardware strategies are both adopted and we started making the hypothesis that the software-implemented self-checking block signature is redundant.

In order to confirm this assumption, we developed a new benchmark implementation where partial hardware is exploited in combination with partial software implementation (sw data+branch+partial hw version), where only data processing instruction replication and conditional branch instruction duplication are exploited. When performing the injection of 1,000 randomly selected faults for each of the pipeline stages of the considered processor, we observed the results reported in Tables 5 and 6.

It can be observed from the reported results that the nonredundant version achieves a complete fault coverage and the software technique for the control flow monitoring can be successfully removed once the hardware has been replicated. Indeed, the function of identifying the occurrence of erroneous branches is already performed by partial hardware replication and its replica becomes useless. The simplification of the adopted software hardening has positive effects on the overheads the method introduces, in particular on the execution time. We indeed observed a reduction of the time overhead ranging from a modest 3.5% (EWF) of the simpler benchmark up to 19.8% of the most complex one (KLMN). Since the self-checking block signature technique has to be inserted in any basic block, the overhead it introduces increases with the complexity

Table 6 Overheads for the nonredundant sw/hw implementation with a complete fault coverage

Program	Version	Exec time ratio	Data size ratio	Code size ratio	Memory ratio
FIR	data+br+hw	3.94	2.00	2.86	2.69
	data+hw	3.44	2.00	2.57	2.46
EWF	data+br+hw	10.86	2.00	5.89	5.17
	data+hw	10.65	2.00	5.65	4.97
KLMN	data+br+hw	6.32	2.00	4.01	2.32
	data+hw	4.59	2.00	3.14	2.18

of the benchmark. When complex programs have to be hardened, its implementation through dedicated hardware seems thus the best option.

In our experiments we also investigated the possibility of further reducing the amount of adopted software techniques. We indeed produced an alternative implementation where only data processing instruction replication and partial hardware redundancy are used (sw data+partial hw implementation). By further reducing the amount of additional instructions the processor has to execute, we further reduce the time overhead (as well as the code overhead) the technique introduces. However, the experiments we performed by injecting 1,000 randomly selected faults for each pipeline stage showed a not complete fault detection, indeed few faults exist that escape the implemented hardening strategies. These faults alter the execution of data-processing instructions whose results affect the program execution flow (e.g., the instructions used to compute the Boolean function used as a condition in a loop). For coping with these faults conditional branch instructions must be duplicated, or different hardware redundancy techniques have to be introduced.

The adoption of these software fault injection campaigns allows us to verify the validity of the proposed analysis; radiation ground testing experiments could better support our claims, although, as it has been shown in [12], the adopted simulation framework matches the radiation testing experiments.

4 Conclusion

The proposed solution is aimed at covering with hardware techniques the faults not addressed by means of software techniques for fault detection. Previous approaches have been refined, from the methodological point of view, thus allowing an efficient application of software and hardware techniques for fault detection, in terms of overheads and benefits. The result is a combined set of techniques, that maintains the flexibility of the software, as well as the possibility to introduce such redundancy only for critical tasks, and introduces the minimal amount of hardware redundancy and processor modification, needed to provide a complete fault coverage, when desired. The proposed solution, derived from the qualitative and systematic analysis of fault/error relations, carried out from an architectural independent point of view, has driven the modification of the processor architecture to achieve a complete fault coverage. Two fault injection campaigns have been carried out on a specific architecture, to evaluate the coverage of both hardware and software techniques, and the experimental

results confirmed the theoretical methodological, qualitative analysis.

References

1. Alkhalifa Z, Nair VSS, Krishnamurthy N, Abraham JA (1999) Design and evaluation of system-level checks for on-line control flow error detection. *IEEE Trans Parallel Distrib Syst* 10(6):627–641
2. Bernardi P, Bolzani L, Rebaudengo M, Sonza Reorda M, Violante M (2005) An integrated approach for increasing the soft-error detection capabilities in SoCs processors. In *Proc. 20th IEEE Int. Symp. on Defect and Fault Tolerance in VLSI Systems*, pp 445–453
3. Bolchini C, Miele A, Pomante L, Salice F, Sciuto D (2004) Reliable system co-design: the FIR case study. In *Proc. 19th IEEE Int. Symp. on Defect and Fault Tolerance in VLSI Systems*. Cannes, France, pp 433–441
4. Bolchini C, Miele A, Salice F, Sciuto D (2005) A model of soft error effects in generic IP processors. In *Proc. 20th IEEE Int. Symp. on Defect and Fault Tolerance in VLSI Systems*, pp 334–342
5. Bolchini C, Pomante L, Salice F, Sciuto D (2005) Reliable system specification for self-checking data-paths. In *Proc. of the Conf. on Design, Automation and Test in Europe*. IEEE Computer Society, Washington, DC, USA, pp 1278–1283
6. Bolchini C, Miele A, Rebaudengo M, Sterpone L, Violante M, Sciuto D (2006) Combined software and hardware techniques for the design of reliable IP processors. In *Proc. IEEE Int. Symp. on Defect and Fault Tolerance in VLSI Systems*, pp 265–273
7. Civera P, Macchiarulo L, Rebaudengo M, Sonza Reorda M, Violante M (2002) An FPGA-based approach for speeding-up fault injection campaigns on safety-critical circuits. *Journal of Electronic Testing: Theory and Applications* 18(3):261–271, June
8. Dodd PE, Massengill LW (2003) Basic mechanisms and modeling of single-event upset in digital microelectronics. *IEEE Trans Nucl Sci* 50(3):583–602, June
9. Gaisler J (2003) The LEON2 IEEE-1754 (SPARC V8) Processor <http://www.gaisler.com>
10. Goloubeva O, Rebaudengo M, Reorda MS, Violante M (2005) Improved software-based processor control-flow errors detection technique. In *Reliability and Maintainability Symposium, 2005. Proceedings. Annual*, pp 583–589
11. Rebaudengo M, Sonza Reorda M, Torchiano M, Violante M (1999) Soft-error detection through software fault-tolerance techniques. In *Proc. IEEE Int. Symp. on Defect and Fault Tolerance in VLSI Systems*, pp 210–218
12. Rebaudengo M, Sonza Reorda M, Violante M, Nicolescu B, Velazco R (2002) Coping with SEUs/SETs in microprocessors by means of low-cost solutions: a comparative study. *IEEE Trans Nucl Sci* (49)3:1491–1495, June
13. Ziegler F et al (1996) Terrestrial cosmic rays and soft errors. *IBM J Res Develop* 40(1)

Cristiana Bolchini is an Associate Professor at Politecnico di Milano. She graduated in 1993 (laurea degree) in Electronic Engineering from

the same institution. She joined the Computer Architecture group working on reliable design methodologies and in 1997 she received her Ph.D. in Computer Science and Automation Engineering from the Dipartimento di Elettronica e Informazione of the Politecnico di Milano. Dr. Bolchini is an associate editor of the IEEE Transactions on Computers and participates to the Technical Program Committee of conferences and symposia in the area of test and fault tolerance for digital systems and is a reviewer for journals and transactions in this same area. Her research interests are digital system design with a specific focus on reliability properties, hardware/software co-design of dependable systems and reconfigurable systems. She has authored several papers in this area.

Antonio Miele graduated in Computer Science Engineering from Politecnico di Milano in 2006, and received a MS degree in Computer Science from the University of Illinois at Chicago. He is currently a Ph.D. student in Information Technology at Politecnico di Milano. His research interests include reliable design methodologies for embedded systems and hardware/software co-design.

Maurizio Rebaudengo received the M.S. degree in Electronic Engineering (1991), and the Ph.D. degree in Computer Engineering (1995), both from Politecnico di Torino, Torino, Italy. Currently, he is an Associate Professor at the Department of Computer Engineering in the same Institution. His research interests include Testing and Dependability analysis of computer-based systems.

Fabio Salice received his PhD in Electronics and Telecommunication in 1995 from the Politecnico di Milano where, currently, he covers the position of Associate Professor at the Dipartimento di Elettronica e Informazione in the Como Campus. He also covers the position of Senior Specialist consultant at CEFRIEL in the area of Embedded System Design. Fabio is member IEEE and he is a reviewer for journals and transactions in the area of electronic design automation and test/fault tolerance for digital systems. His research interests are digital system design with a specific focus on reliability properties,

hardware/software co-design of dependable systems and ESL tools and methods. He has authored several papers and he has received 3 awards: two best posters and one best paper.

Donatella Sciuto received her Laurea in Electronic Engineering from Politecnico di Milano and her PhD in Electrical and Computer Engineering from the University of Colorado, Boulder. She is currently a Full Professor at the Dipartimento di Elettronica e Informazione of the Politecnico di Milano, Italy. She is member IEEE, IFIP 10.5, EDAA. She is or has been member of different program committees of EDA conferences: DAC, ICCAD, DATE, CODES+HSSS, DFT, FDL, and associate Editor of the IEEE Transactions on Computers, IEEE and the Journal of Design Automation of Embedded Systems, Kluwer Academic Publishers. She is in the executive committee of the conference IEEE/ACM Design Automation and Test in Europe, for which she has been Program Chair in 2006, Vice General Chair for 2007 and is General Chair in 2008. Her main research interests cover the methodologies for hardware/software co-design of embedded systems, from the specification level down to the implementation of both the hardware and software components, including reconfigurable systems.

Luca Sterpone received the M.S. degree in Computer Engineering in 2003 and the Ph.D. in computer Engineering in 2007, both from Politecnico di Torino, Torino, Italy. Currently, he is an Assistant Researcher at the Department of Computer Engineering in the same Institution. His research interests include fault tolerant systems, place and route algorithms and reconfigurable computing.

Massimo Violante received the M.S. degree in Computer Engineering (1996) and the Ph.D. degree in Computer Engineering (2001) from Politecnico di Torino, Italy. Currently, he is an Assistant Professor with the Department of Computer Engineering of the same Institution. His main research interests include design, validation, and test of fault-tolerant electronic systems with particular emphasis on reconfigurable systems. Dr. Violante published more than 120 papers on these topics, including a book on Software-Implemented Hardware Fault Tolerance.